

Problem Statement: Which model is suitable for given dataset

Importing Packages

In [3]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

Data Collection

In [4]:

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\insurance.csv")
df
```

Out[4]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

Data Cleaning and Preprocessing

In [5]:

```
df.head()
```

Out[5]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [6]:

```
df.tail()
```

Out[6]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [8]:

```
df.describe()
```

Out[8]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [9]:

```
df.columns
```

Out[9]:

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dt  
type='object')
```

In [10]:

```
df.shape
```

Out[10]:

```
(1338, 7)
```

Duplicate Values

In [11]:

```
df.isnull().sum()
```

Out[11]:

```
age      0  
sex      0  
bmi      0  
children 0  
smoker   0  
region   0  
charges  0  
dtype: int64
```

Data Visualization

In [12]:

```
df['smoker'].value_counts()
```

Out[12]:

```
smoker
no      1064
yes      274
Name: count, dtype: int64
```

In [13]:

```
df['sex'].value_counts()
```

Out[13]:

```
sex
male      676
female    662
Name: count, dtype: int64
```

In [14]:

```
df['region'].value_counts()
```

Out[14]:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [15]:

```
s=pd.crosstab(df['smoker'],df['sex'])
print(s)
```


```
sex      female  male
smoker
no          547   517
yes          115   159
```

In [16]:

```
s.plot(kind='bar', stacked=True, color=['lightcoral', 'LightSeaGreen'], grid=False)
```

Out[16]:

<Axes: xlabel='smoker'>

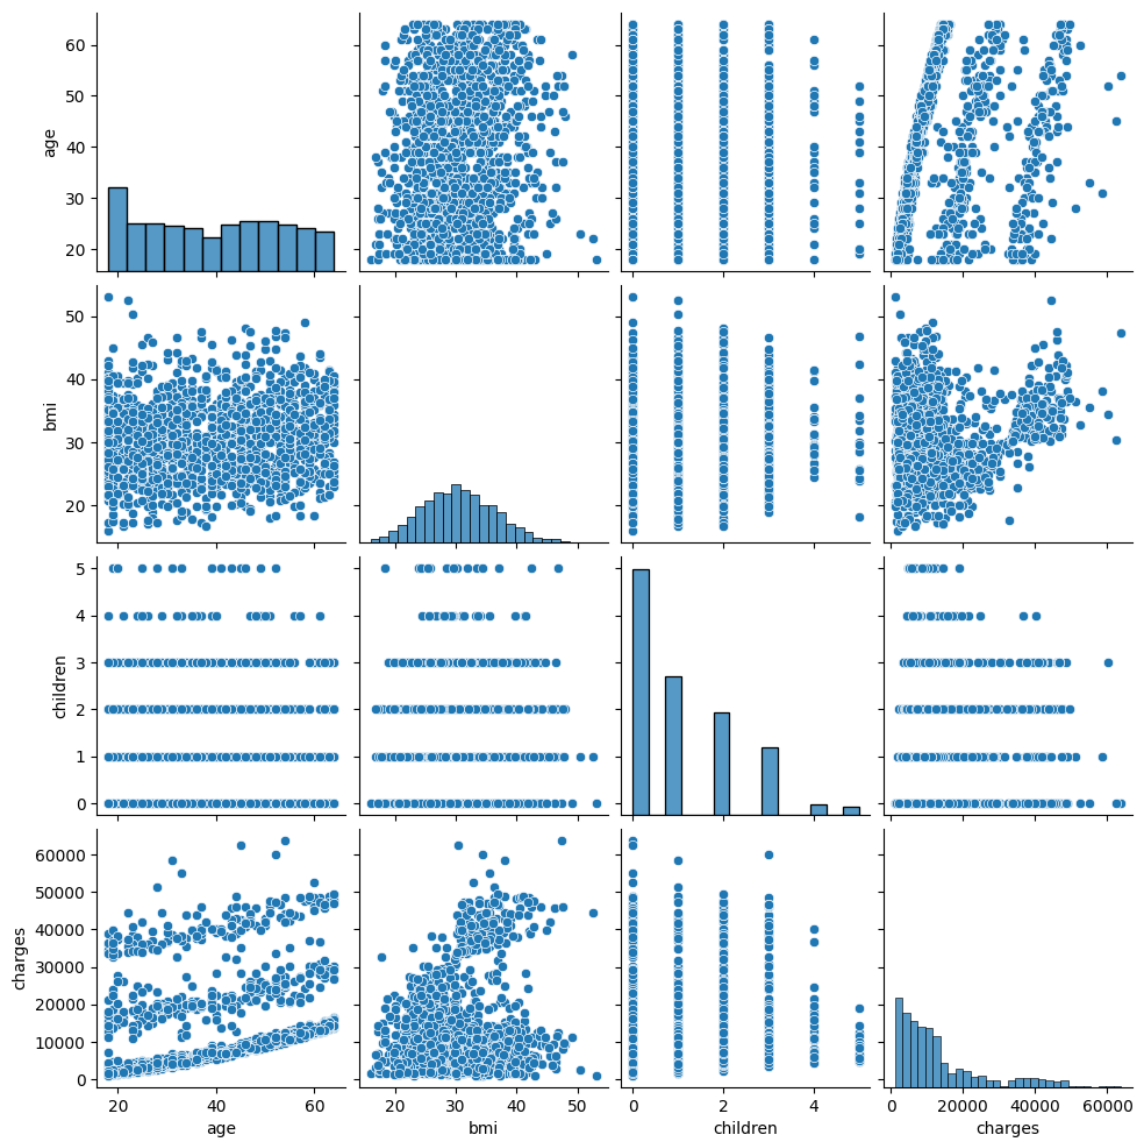


In [17]:

```
sns.pairplot(df)
```

Out[17]:

<seaborn.axisgrid.PairGrid at 0x1f9fc738f40>



In [18]:

```
c=pd.crosstab(df['age'],df['sex'])  
print(c)
```


sex	female	male
age		
18	33	36
19	33	35
20	14	15
21	13	15
22	13	15
23	14	14
24	14	14
25	13	15
26	13	15
27	14	14
28	14	14
29	13	14
30	13	14
31	13	14
32	13	13
33	13	13
34	13	13
35	12	13
36	12	13
37	12	13
38	13	12
39	13	12
40	13	14
41	13	14
42	13	14
43	14	13
44	14	13
45	14	15
46	14	15
47	15	14
48	15	14
49	14	14
50	14	15
51	15	14
52	15	14
53	14	14
54	14	14
55	13	13
56	13	13
57	13	13
58	13	12
59	13	12
60	11	12
61	12	11
62	12	11
63	12	11
64	11	11

In [19]:

```
c.plot(kind='line', stacked=False, color=['aquamarine','lightcoral'],grid=False)
```

Out[19]:

<Axes: xlabel='age'>



In [20]:

```
s = {'region':{'northeast':1,'northwest':2,'southwest':3,'southeast':4}}
df = df.replace(s)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	3	16884.92400
1	18	male	33.770	1	no	4	1725.55230
2	28	male	33.000	3	no	4	4449.46200
3	33	male	22.705	0	no	2	21984.47061
4	32	male	28.880	0	no	2	3866.85520
...
1333	50	male	30.970	3	no	2	10600.54830
1334	18	female	31.920	0	no	1	2205.98080
1335	18	female	36.850	0	no	4	1629.83350
1336	21	female	25.800	0	no	3	2007.94500
1337	61	female	29.070	0	yes	2	29141.36030

[1338 rows x 7 columns]

In [21]:

```
S = {'sex':{'female':1,'male':2}}
df =df.replace(S)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	3	16884.92400
1	18	2	33.770	1	no	4	1725.55230
2	28	2	33.000	3	no	4	4449.46200
3	33	2	22.705	0	no	2	21984.47061
4	32	2	28.880	0	no	2	3866.85520
...
1333	50	2	30.970	3	no	2	10600.54830
1334	18	1	31.920	0	no	1	2205.98080
1335	18	1	36.850	0	no	4	1629.83350
1336	21	1	25.800	0	no	3	2007.94500
1337	61	1	29.070	0	yes	2	29141.36030

[1338 rows x 7 columns]

In [22]:

```
x = df.drop('smoker',axis=1)
y = df['smoker']
```

Linear Regression

In [25]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=40)
x_train.shape,x_test.shape
```

Out[25]:

((936, 6), (402, 6))

Logistic Regression

In [27]:

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

0.9378109452736318

Decision Tree

In [28]:

```
from sklearn.tree import DecisionTreeClassifier
clf = DecisionTreeClassifier(random_state=0)
clf.fit(x_train,y_train)
score = clf.score(x_test,y_test)
print(score)
```

0.9601990049751243

RandomForestClassifier

In [29]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
print(rfc.score(x_test,y_test))
```

0.9626865671641791


In [30]:

```
params={'max_depth':[2,5,10,20,25], 'min_samples_leaf':[5,20,30,50,100,200], 'n_estimators
```

In [31]:

```
from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

Out[31]:



In [32]:

```
grid_search.best_score_
```

Out[32]:

0.9497863247863247

In [33]:

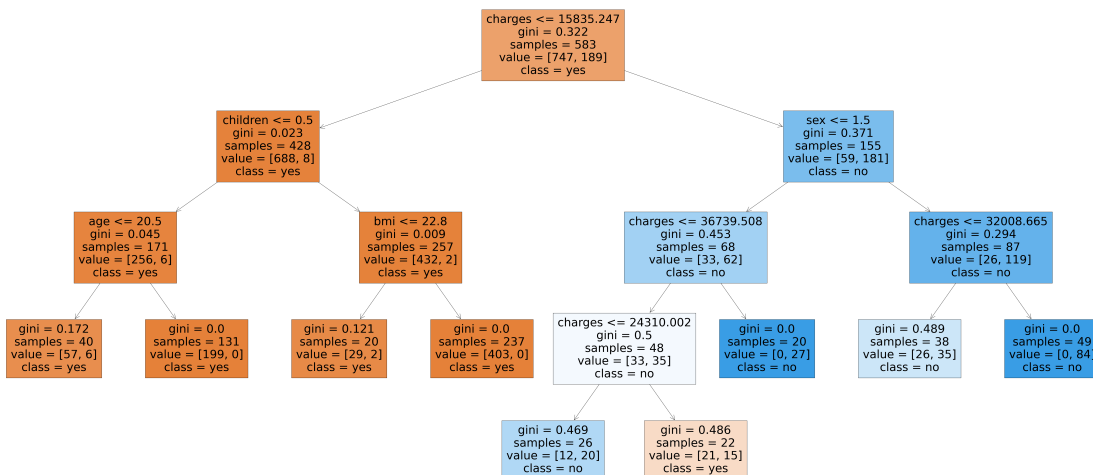
```
rfc_best = grid_search.best_estimator_
```

In [34]:

```
from sklearn.tree import plot_tree
plt.figure(figsize = (90,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['yes','no'],fille
```

Out[34]:

```
[Text(0.5, 0.9, 'charges <= 15835.247\ngini = 0.322\nsamples = 583\nnvalue = [747, 189]\nnclass = yes'),
 Text(0.25, 0.7, 'children <= 0.5\ngini = 0.023\nsamples = 428\nnvalue = [688, 8]\nnclass = yes'),
 Text(0.125, 0.5, 'age <= 20.5\ngini = 0.045\nsamples = 171\nnvalue = [256, 6]\nnclass = yes'),
 Text(0.0625, 0.3, 'gini = 0.172\nsamples = 40\nnvalue = [57, 6]\nnclass = yes'),
 Text(0.1875, 0.3, 'gini = 0.0\nsamples = 131\nnvalue = [199, 0]\nnclass = yes'),
 Text(0.375, 0.5, 'bmi <= 22.8\ngini = 0.009\nsamples = 257\nnvalue = [432, 2]\nnclass = yes'),
 Text(0.3125, 0.3, 'gini = 0.121\nsamples = 20\nnvalue = [29, 2]\nnclass = yes'),
 Text(0.4375, 0.3, 'gini = 0.0\nsamples = 237\nnvalue = [403, 0]\nnclass = yes'),
 Text(0.75, 0.7, 'sex <= 1.5\ngini = 0.371\nsamples = 155\nnvalue = [59, 181]\nnclass = no'),
 Text(0.625, 0.5, 'charges <= 36739.508\ngini = 0.453\nsamples = 68\nnvalue = [33, 62]\nnclass = no'),
 Text(0.5625, 0.3, 'charges <= 24310.002\ngini = 0.5\nsamples = 48\nnvalue = [33, 35]\nnclass = no'),
 Text(0.5, 0.1, 'gini = 0.469\nsamples = 26\nnvalue = [12, 20]\nnclass = no'),
 Text(0.625, 0.1, 'gini = 0.486\nsamples = 22\nnvalue = [21, 15]\nnclass = yes'),
 Text(0.6875, 0.3, 'gini = 0.0\nsamples = 20\nnvalue = [0, 27]\nnclass = no'),
 Text(0.875, 0.5, 'charges <= 32008.665\ngini = 0.294\nsamples = 87\nnvalue = [26, 119]\nnclass = no'),
 Text(0.8125, 0.3, 'gini = 0.489\nsamples = 38\nnvalue = [26, 35]\nnclass = no'),
 Text(0.9375, 0.3, 'gini = 0.0\nsamples = 49\nnvalue = [0, 84]\nnclass = no')]
```

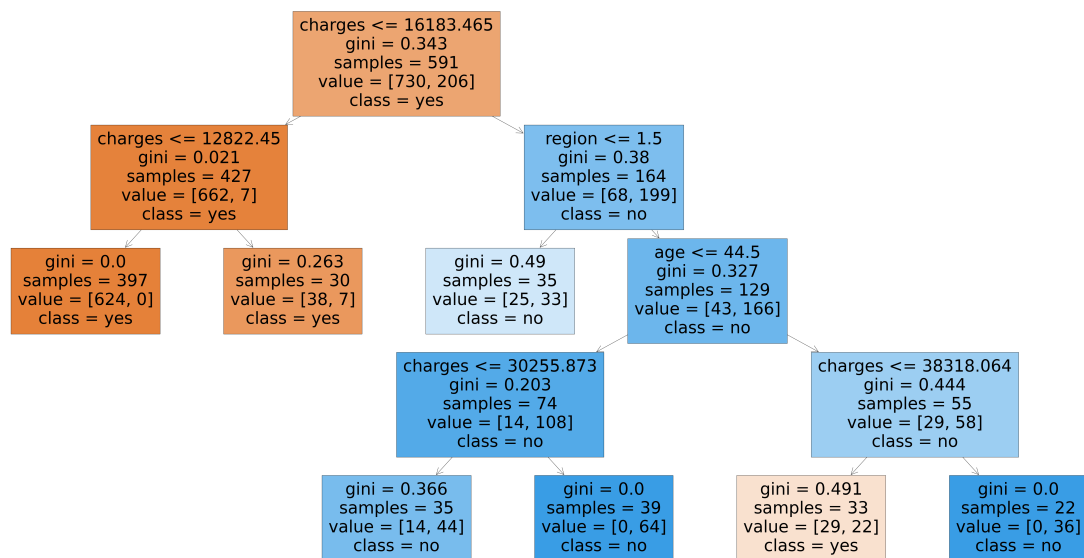


In [35]:

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[7],feature_names=x.columns,class_names=['yes','no'],fille
```

Out[35]:

```
[Text(0.36363636363636365, 0.9, 'charges <= 16183.465\ngini = 0.343\nsampl
es = 591\nvalue = [730, 206]\nnclass = yes'),
 Text(0.18181818181818182, 0.7, 'charges <= 12822.45\ngini = 0.021\nsampl
es = 427\nvalue = [662, 7]\nnclass = yes'),
 Text(0.09090909090909091, 0.5, 'gini = 0.0\nsamples = 397\nvalue = [624,
0]\nnclass = yes'),
 Text(0.2727272727272727, 0.5, 'gini = 0.263\nsamples = 30\nvalue = [38,
7]\nnclass = yes'),
 Text(0.5454545454545454, 0.7, 'region <= 1.5\ngini = 0.38\nsamples = 164
\nvalue = [68, 199]\nnclass = no'),
 Text(0.45454545454545453, 0.5, 'gini = 0.49\nsamples = 35\nvalue = [25, 3
3]\nnclass = no'),
 Text(0.6363636363636364, 0.5, 'age <= 44.5\ngini = 0.327\nsamples = 129\n
value = [43, 166]\nnclass = no'),
 Text(0.45454545454545453, 0.3, 'charges <= 30255.873\ngini = 0.203\nsampl
es = 74\nvalue = [14, 108]\nnclass = no'),
 Text(0.36363636363636365, 0.1, 'gini = 0.366\nsamples = 35\nvalue = [14,
44]\nnclass = no'),
 Text(0.5454545454545454, 0.1, 'gini = 0.0\nsamples = 39\nvalue = [0, 64]
\nnclass = no'),
 Text(0.8181818181818182, 0.3, 'charges <= 38318.064\ngini = 0.444\nsampl
es = 55\nvalue = [29, 58]\nnclass = no'),
 Text(0.7272727272727273, 0.1, 'gini = 0.491\nsamples = 33\nvalue = [29, 2
2]\nnclass = yes'),
 Text(0.9090909090909091, 0.1, 'gini = 0.0\nsamples = 22\nvalue = [0, 36]
\nnclass = no')]
```



In [36]:

```
rfc_best.feature_importances_
```

Out[36]:

```
array([0.02957835, 0.00478561, 0.02999727, 0.00681847, 0.00549181,
       0.92332849])
```

In [37]:

```
imp_df = pd.DataFrame({"Varname":x_train.columns,'Imp':rfc_best.feature_importances_})
imp_df.sort_values(by='Imp',ascending=False)
```

Out[37]:

	Varname	Imp
5	charges	0.923328
2	bmi	0.029997
0	age	0.029578
3	children	0.006818
4	region	0.005492
1	sex	0.004786

In [38]:

```
df['bmi'].value_counts()
```

Out[38]:

```
bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
..
46.200     1
23.800     1
44.770     1
32.120     1
30.970     1
Name: count, Length: 548, dtype: int64
```

conclusion:-

Based dataset We conclude that male smoker are high compared to female smokers.we conclude that "Logistic regression" is the best model for the given data set

In [39]:

```
import pickle
```

In [40]:

```
filename="Insurance prediction"
pickle.dump(lr,open(filename,'wb'))
```

In []: